

Assignment No 6: Object detection using Transfer Learning of CNN architectures

- Load in a pre-trained CNN model trained on a large dataset
- Freeze parameters (weights) in model's lower convolutional layers
- Add custom classifier with several layers of trainable parameters to model
- Train classifier layers on training data available for task
- Fine-tune hyper parameters and unfreeze more layers as needed

```
In [1]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
In [2]: train_dir = 'dataset/mnist-jpg/mnist-jpg/train/'  
test_dir = 'dataset/mnist-jpg/mnist-jpg/test/'
```

```
In [13]: img_gen = ImageDataGenerator(rescale=1.0/255)  
  
data_gen = img_gen.flow_from_directory(  
    train_dir,  
    target_size=(32,32),  
    batch_size=5000,  
    shuffle=True,  
    class_mode='categorical'  
)
```

Found 60000 images belonging to 10 classes.
Found 60000 images belonging to 10 classes.

```
In [14]: x_train, y_train = data_gen[0]  
x_test, y_test = data_gen[2]
```

```
In [15]: from tensorflow.keras.applications import VGG16  
path = 'dataset/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5'  
  
vgg_model = VGG16(weights=path, include_top=False, input_shape=(32,32,3))
```

```
In [16]: for layer in vgg_model.layers:  
    layer.trainable=False
```

```
In [17]: from tensorflow import keras  
from tensorflow.keras.layers import Dense, Flatten, Dropout
```

```
In [22]: custom_classifier = keras.Sequential([  
    Flatten(input_shape=(1,1,512)),  
    Dense(100, activation='relu'),  
    Dropout(0.2),  
    Dense(100, activation='relu'),  
    Dropout(0.2),  
    Dense(10, activation='softmax')  
)  
  
model = keras.Sequential([  
    vgg_model,  
    custom_classifier  
)
```

```
In [23]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [25]: model.fit(x_train, y_train, batch_size=100, epochs=1, validation_data=(x_test,y_test))
```

50/50 [=====] - 204s 4s/step - loss: 1.6244 - accuracy: 0.3820 - val_loss: 1.2011 - val_accuracy: 0.5770

```
Out[25]: <keras.src.callbacks.History at 0x22c7ef89c60>
```

```
In [26]: for layer in vgg_model.layers[:-4]:  
    layer.trainable = True  
  
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
model.fit(x_train, y_train, batch_size=1000, epochs=1, validation_data=(x_test,y_test))
```

5/5 [=====] - 208s 42s/step - loss: 2.5760 - accuracy: 0.3306 - val_loss: 2.1499 - val_accuracy: 0.2630

```
Out[26]: <keras.src.callbacks.History at 0x22c7822e950>
```

```
In [27]: loss, acc = model.evaluate(x_test, y_test)
```

```
print(loss, " ", acc)
```

```
32/32 [=====] - 4s 120ms/step - loss: 2.1499 - accuracy: 0.2630  
2.1498725414276123 0.2630000114440918
```


```
In [28]: pred = model.predict(x_test)
```

```
32/32 [=====] - 4s 115ms/step
```

```
In [29]: labels = list(data_gen.class_indices.keys())
```

```
In [32]: import matplotlib.pyplot as plt  
import numpy as np  
plt.imshow(x_test[10])  
plt.title(str(labels[np.argmax(pred[10])]))  
print(str(labels[np.argmax(y_test[10])]))
```

0



```
In [33]: y_test[10]
```

```
Out[33]: array([1., 0., 0., 0., 0., 0., 0., 0., 0., 0.], dtype=float32)
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js